

Lecture 5


Local Search

5.1 Local Search Algorithms

In many search problems, the path from the initial state to the goal is **not** (4.1) important, but the final configuration of the goal is the target of the problem, e.g., 8-queens problem.

Local search is a group of algorithms that focus on this kind of problems. Since the path is not matter, local search algorithms typically keep only one node while operating.

Apart from searching for goal states, the algorithms can be used to solve **optimization problems**. The optimization problems aim to find the state with the best value of an **objective function**.



5.2 Hill Climbing Search

Hill climbing search algorithm repeatedly moves the current node to the direction that increases (or decreases) the objective value. It ends when the current node reaches a peak where no neighbor has a higher value. (4.1.1)

Code 5.1 Hill Climbing Algorithm implemented in Python


```
1 def hill_climbing(s0, succ, f):  
2     # s0 = initial state  
3     # succ = successor function  
4     # f = evaluation function  
5     u = s0  
6     while True:  
7         s = succ(u)  
8         v = max(s, key=f)  
9         if f(v) > f(u):  
10            u = v  
11        else:  
12            return u
```

Example 5.1 To solve the 8-queens problem by the Hill Climbing algorithm, we need to define a function that evaluates each state.

We can use the number of attacks among the queens on the board as an evaluation value of a state. Thus, $f(s)$ returns the number of attacks given a state s . The search can start from a randomly created board with one queen placed in each column.

A successor state can be generated by selecting one queen and change its row position.

Here, we search for the state with $f = 0$.



Exercise 5.1 A local search problem represents a state by an ordered pair. The initial state is $(0,0)$. The successors of a state (x,y) are $(x-1,y)$, $(x+1,y)$, $(x,y-1)$, and $(x,y+1)$. The objective function is shown in the following figure. Use the hill climbing algorithm to search for the state with maximum objective value.

	0	1	2	3	4	5	6	7	8	9
0	3	4	4	4	5	4	5	5	5	4
1	4	5	5	5	4	4	5	6	5	3
2	5	6	7	6	4	3	6	7	6	2
3	5	6	6	6	4	2	6	6	6	2
4	5	5	5	4	4	1	6	5	5	3
5	3	3	4	3	7	7	7	6	4	4
6	4	3	4	3	7	8	7	8	7	5
7	4	5	5	4	6	7	9	8	7	5
8	4	5	5	4	6	8	8	8	7	6
9	3	3	3	4	6	6	7	7	7	5

5.2.1 Random-restart Hill Climbing

To avoid getting stuck in local maxima, the hill climbing search can be conducted multiple times. In each iteration, it starts from different randomly generated initial states.

Exercise 5.2 Conduct the random-restart hill climbing on the following space.

	0	1	2	3	4	5	6	7	8	9
0	3	4	4	4	5	4	5	5	5	4
1	4	5	5	5	4	4	5	6	5	3
2	5	6	7	6	4	3	6	7	6	2
3	5	6	6	6	4	2	6	6	6	2
4	5	5	5	4	4	1	6	5	5	3
5	3	3	4	3	7	7	7	6	4	4
6	4	3	4	3	7	8	7	8	7	5
7	4	5	5	4	6	7	9	8	7	5
8	4	5	5	4	6	8	8	8	7	6
9	3	3	3	4	6	6	7	7	7	5

Code 5.2 Random Restart Hill Climbing Algorithm

```
1 from hillclimbing import *
2
3
4 def random_restart_hill_climbing(succ, f, rnd, n):
5     # succ = successor function
6     # f     = evaluation function
7     # rnd   = function that randomly generates state
8     # n     = number of iterations
9     best_value = -float("inf")
10    for i in range(n):
11        s0 = rnd()
12        u = hill_climbing(s0, succ, f)
13        if f(u) > best_value:
14            best_state = u
15            best_value = f(u)
16    return best_state
```

Handwritten notes:
→ defined by user (pointing to n)
→ start from a random state (pointing to s0)

Random-restart hill climbing provides a higher chance to get a better solution.

The maximum is NOT guaranteed by this algorithm.

5.2.2 Stochastic Hill Climbing Search

The hill climbing search basically goes directly to the closest local maximum ~~with~~ exploring the search space.

with
out

To allow the search to explore the space, we use the probability p to judge if a newly generated state should be accepted.

after generating a set of successors, this algorithm just randomly selects one successor.
Calculate probability p

$$p = \frac{1}{1 + e^{\frac{\text{current} - \text{new}}{T}}}$$

where T is the parameter to control how easy the worse state is selected.

Code 5.3 Stochastic Hill Climbing Algorithm

```


1 import random
2 import math
3
4
5 def stochastic_hill_climbing(s0, succ, f, n, t):
6     # succ = successor function
7     # f = evaluation function
8     # rnd = function that randomly generates state
9     # n = number of iterations
10    best_value = -float("inf")
11    u = s0
12    for i in range(n):
13        s = succ(u) = generate a set of successors
14        v = random.choice(s) = randomly pick one successor
15        p = 1/(1+math.exp((f(u)-f(v))/t)) = calculate p
16        r = random.uniform(0.0, 1.0) = pick a random value
17        if r < p:
18            u = v
19        if f(u) > best_value:
20            best_state = u
21            best_value = f(u)
22    return best_state

```

P
current ? new

u = current
v = new

r (move to "new")



Consider the following values of p based on the different values of T

new is better than current new is worse than current

T	$e^{-\frac{13}{T}}$	p	$e^{\frac{13}{T}}$	p
1	0.00000226	1.000	442413.39	0.00000226
5	0.0498	0.953	13.5	0.0691
10	0.273	0.786	3.67	0.214
20	0.522	0.657	1.92	0.343
50	0.771	0.565	1.30	0.435
10^{10}	1.00	0.500	1.00	0.500

5.3 Simulated Annealing

Simulated annealing is based on an idea from thermodynamics.

- To grow a crystal, we heat the materials to a molten state. Then, cool it down until the crystal structure is frozen in.
- If the cooling is done too quickly, the materials become brittle.

The parameter T in the stochastic hill-climbing can be compared to the temperature. We start with T in a high value to allow random walks, and we decrease T when the search is going on. This allows the exploration of the search space at the beginning, then it goes to the optimal state in the end.

Simulated annealing = Stochastic HC with varied T

High T --> random search

Low T --> hill climbing

Code 5.4 Simulated Annealing Algorithm

```
1 import random
2 import math
3
4
5 def simulated_annealing(s0, succ, f, t0, cooling):
6     # succ = successor function
7     # f     = evaluation function
8     # rnd   = function that randomly generates state
9     # n     = number of iterations
10    best_value = -float("inf")
11    u = s0
12    t = t0
13    i = 0
14    while t > 0:
15        s = succ(u)
16        v = random.choice(s)
17        if f(u) > f(v):
18            u = v
19        else:
20            p = 1/(1+math.exp((f(u)-f(v))/t))
21            r = random.uniform(0.0, 1.0)
22            if r < p:
23                u = v
24            i += 1
25            t = cooling(t, i)
26            if f(u) > best_value:
27                best_state = u
28                best_value = f(u)
29    return best_state
```

initial value of T (points to `t0`)


function to decrease T (points to `cooling`)

decrease T when the search is going on. (points to `t = cooling(t, i)`)

5.3.1 Cooling Schedule


$$T_i = T_0 - i \frac{T_0 - T_N}{N}$$

initial value
 final value
 # iterations
 i th iteration



$$T_i = T_0 \left(\frac{T_N}{T_0} \right)^{\frac{i}{N}}$$

initial value
 final value (usually $\rightarrow 0$)
 i th iteration
 # iterations



5.4 Local Beam Search

Instead of keeping track of only 1 node, **local beam search** algorithm keeps track of k states while conducting the search. It starts from randomly generated k states. (4.1.3)

The algorithm selects the k best neighbors. It may select all of the neighbors if there are less than k states, and randomly select in the case of ties. Then, these k states are used in the next iteration.

Local beam search with k states is different from conducting k random restarts hill climbing search.

The search stops after (1) a fixed number of iterations,

Exercise 5.3 Conduct the local beam search when the beam size is 3, and the initial states are $(0,0)$; $(1,8)$; $(7,4)$ **(2) no improvement**

	0	1	2	3	4	5	6	7	8	9
0	3	4	4	4	5	4	5	5	5	4
1	4	5	5	5	4	4	5	6	5	3
2	5	6	7	6	4	3	6	7	6	2
3	5	6	6	6	4	2	6	6	6	2
4	5	5	5	4	4	1	6	5	5	3
5	3	3	4	3	7	7	7	6	4	4
6	4	3	4	3	7	8	7	8	7	5
7	4	5	5	4	6	7	9	8	7	5
8	4	5	5	4	6	8	8	8	7	6
9	3	3	3	4	6	6	7	7	7	5

3
5
5

→

6
6
6

→

8
7
7

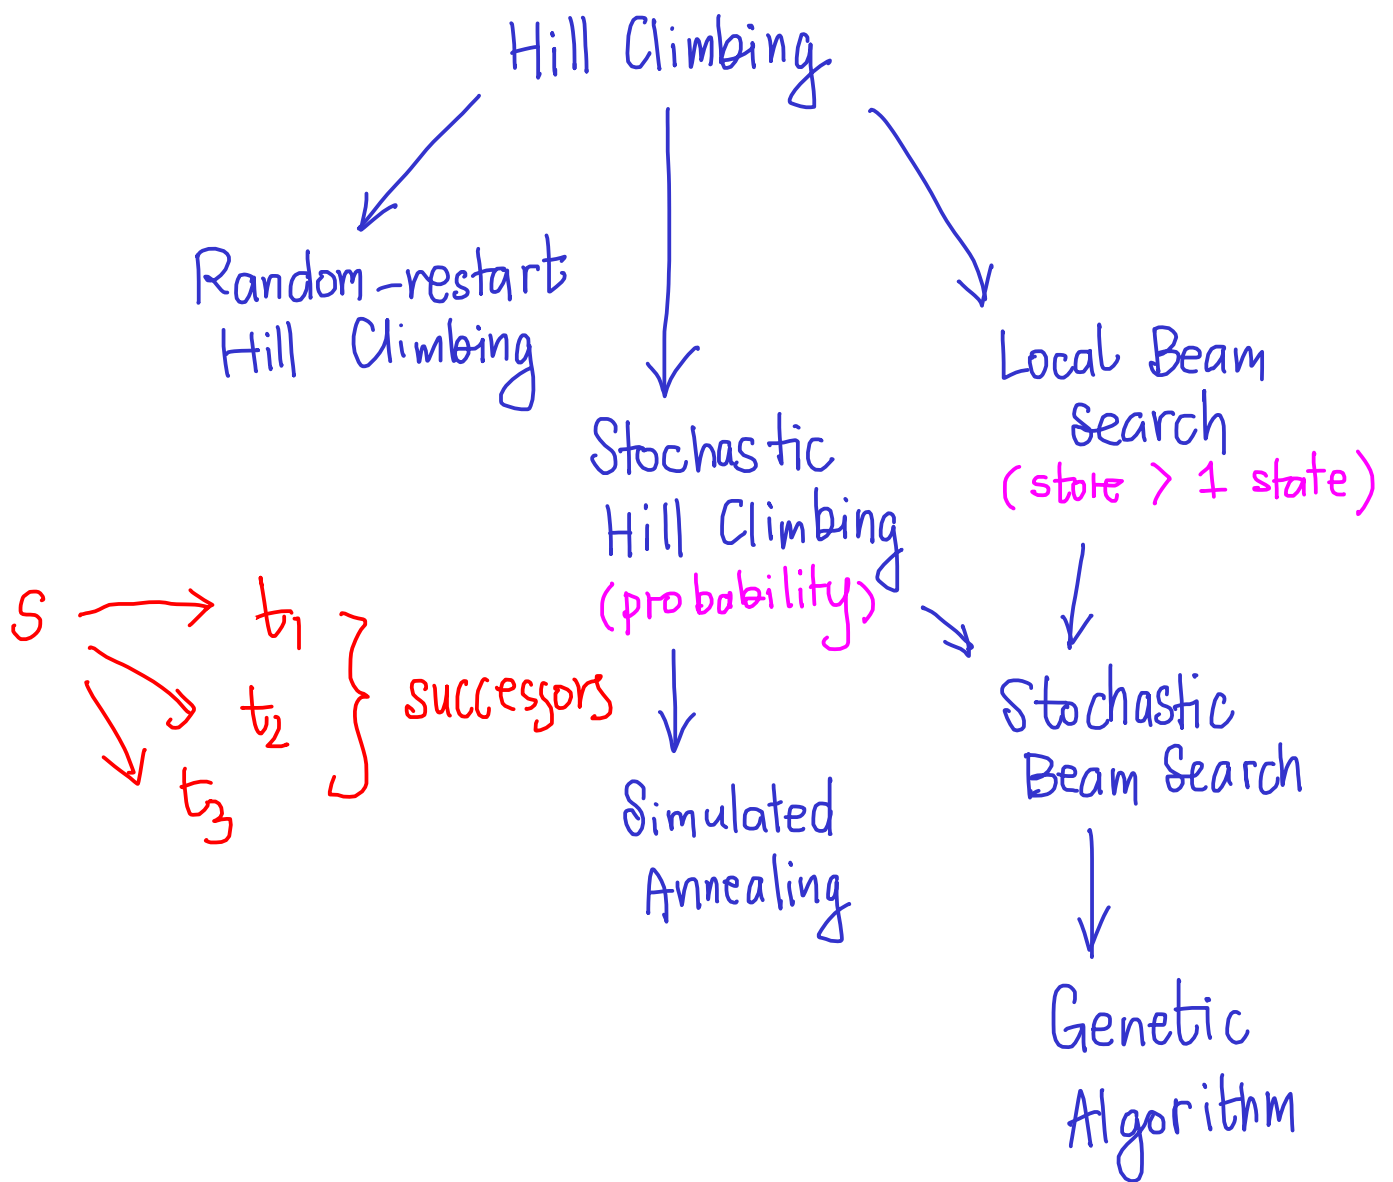
→

8
7
7

no improvement

returns
(7,6)

Local Search



Stochastic beam search chooses k neighbors at random with higher evaluation states are more likely to be selected. Here, we define the probability p as

$$p = e^{-f(s)/T}$$


$f(s)$ = objective value of state s

5.5 Genetic Algorithm

Evolution Theory

Genetic Algorithm is a variant of **stochastic beam search**. Different from (4.1.4) other local search algorithms, GA generates successor states from a pair of states. State = Chromosome

Each state is denoted by a string over a finite set of alphabets, e.g. 0/1. Successor states are generated by two basic operations: crossover, and mutation. This is different from the other local search algorithms.



6 7 2 | 4 6 5 8 8

7 5 2 | 5 1 4 4 7

6 7 2 5 1 4 4 7

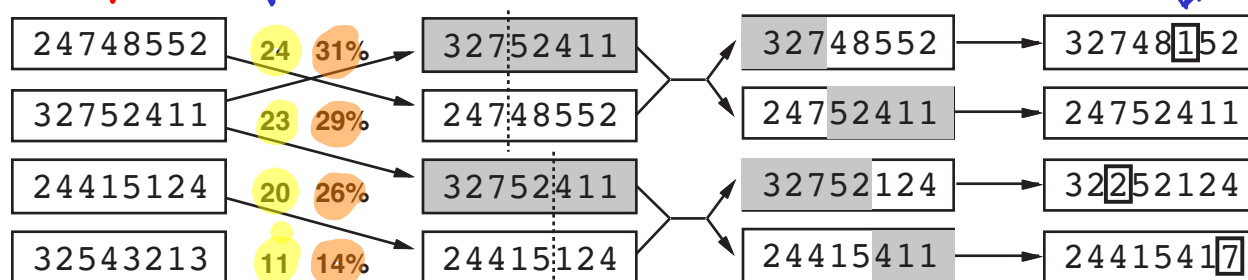
"crossover"

(a successor)

Randomly pick a cutting point

randomly pick value and change it to a random value

Randomly created ↓ Evaluation (higher is better) ↓ value



(a) Initial Population

(b) Fitness Function


(c) Selection

(d) Crossover

(e) Mutation

$$0.14 = \frac{11}{24+23+20+11}$$

probability to be selected



Code 5.5 Genetic Algorithm

```
1 import random, math
2
3 def genetic_algorithm(population, f, n):
4     # population = initial set of individuals
5     # f = evaluation function
6     # n = number of iterations
7     new_population = []
8     for i in range(n):
9         for j in range(len(population)):
10             v = evaluate(population, f)
11             x = select(population, v)
12             y = select(population, v)
13             child = cross_over(x, y)
14             r = random.uniform(0.0, 1.0)
15             if r < 0.05:
16                 child = mutate(child)
17             new_population.append(child)
18         population = new_population
19     return max(population, key=f)
20
21 def evaluate(population, f):
22     return [f(p) for p in population]
23
24 def select(population, val):
25     m = sum(val)
26     r = random.uniform(0, m)
27     c = 0
28     for p, v in zip(population, val):
29         c += v
30         if c > r:
31             return p
32
33 def cross_over(x, y):
34     c = random.randint(0, len(x))
35     return x[0:c] + y[c+1:len(x)]
36
37 def mutate(x):
38     c = random.randint(0, len(x)-1)
39     x[c] = random.randint(0, 1)
40     return x
```

Exercise 5.4 To find the maximum value of the following function

$$f(x, y) = (1 - x^2)e^{-x^2 - y} + x^3 - y^3$$

If x and y are set to vary between -3 and 3 . How can we represent an individual?

(chromosome) find x, y that maximize $f(x, y)$ by using Genetic Algorithm.

an individual = a string over a set of alphabets (e.g. binary string)

a state = a pair of x, y
e.g. $(0.5, -1.28)$

Let us represent an individual by a 16-bit string

10110100 00001011

$x \in [-3, +3]$ $y \in [-3, +3]$

0 = 00000000 $\Rightarrow -3$
 00000001
 $= -3 + \frac{1}{255} = -2.9765$

255 = 11111111 $\Rightarrow +3$

"-3000000 -3000000"

"+3000000 +3000000"

References

Russell, S. and Norvig, P. (2010). Artificial Intelligence: A Modern Approach (3rd edition). Pearson/Prentice Hall.

Michalewicz, F. and Fogel, D. B. (1998). How to Solve It: Modern Heuristics. Springer.