

Lecture 12

Machine Learning

↗ k nearest neighbors
↘ neural networks

12.1 Machine Learning

Machine Learning is a field of study focusing on how to make computer system learn from a set of examples. To be more precise, Machine Learning focuses on automatically constructing models from the collected dataset without requiring explicit programming.

Several types of machine learning techniques have been developed until now, e.g. supervised learning, unsupervised learning, reinforcement learning, etc. In this course, we focus only on supervised machine learning.

(learn from training examples)

12.1.1 Supervised Machine Learning

Given a set of input-output pairs

$$\mathcal{D} = \{(\mathbf{x}_i, t_i)\}_{i=1}^N$$

where,


\mathcal{D} is a training set,

N is the number of examples,

$\mathbf{x}_i \in \mathcal{X}$ is an input vector, and

$t_i \in \mathcal{T}$ is a target output,

find a function $\hat{f}: \mathcal{X} \rightarrow \mathcal{T}$ that predicts the value of t from an unknown \mathbf{x} .



↑ this function is constructed automatically by an ML algorithm.

Each input vector \mathbf{x}_i is usually represented as a D -dimensional feature vector, i.e.

$$\mathbf{x}_i = [x_{i1}, x_{i2}, \dots, x_{iD}]^T \in \mathbb{R}^D$$

← dimension of input vectors

It can be represented as a point in the D -dimensional space.

It is assumed that each pair (\mathbf{x}_i, t_i) is drawn i.i.d. (independently and identically distributed) from an unknown distribution on \mathcal{X} .


12.1.2 Classification and Regression

When \mathcal{T} is a finite set or t_i is categorical, the task of the supervised learning is called classification. The task is called regression when t_i is real-valued i.e. $\mathcal{T} = \mathbb{R}$ or $\mathcal{T} = \mathbb{R}^D$. In this course, we mainly focus on the classification task where the binary classification is the most frequently studied problem. It is the classification task when $\mathcal{T} = \{-1, +1\}$.

→ # possible outputs is finite

→ # possible outputs is infinite

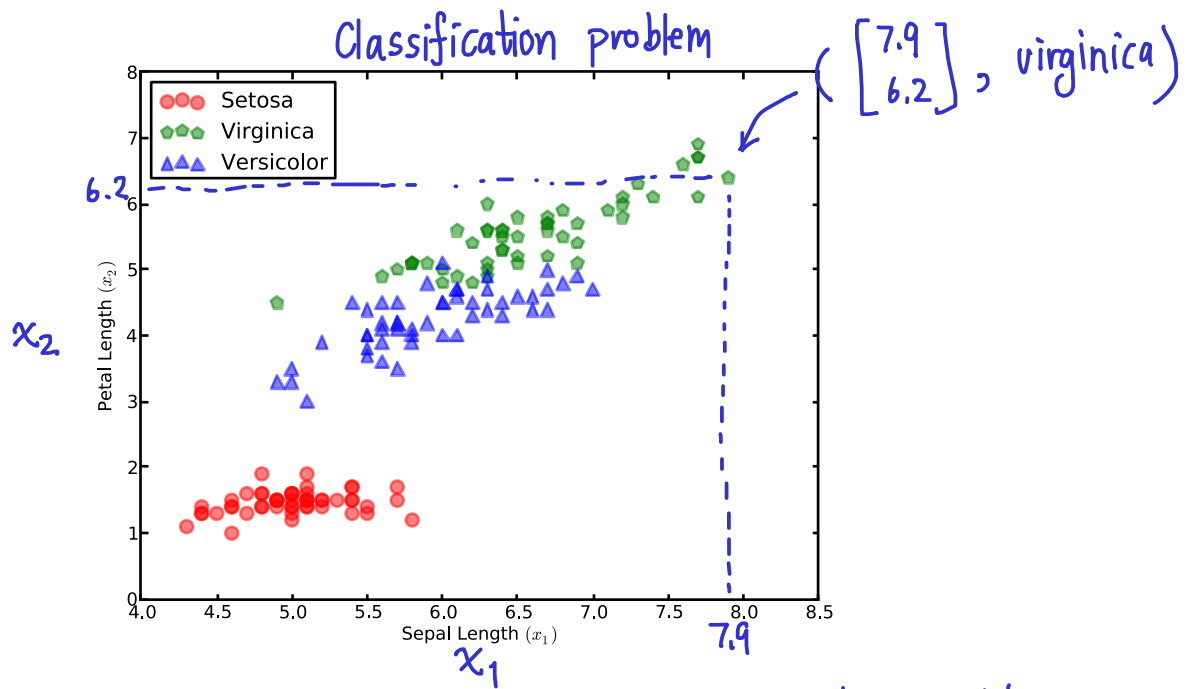
Example 12.1 Predict the housing values in suburbs of Boston.¹



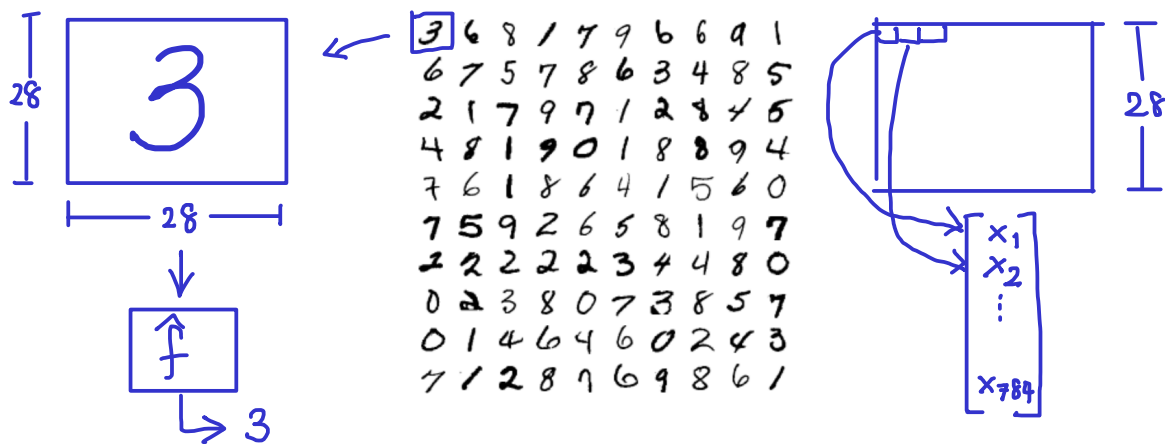
Example 12.2 Predict the class of iris using its shape.²

¹Housing Data Set – <http://archive.ics.uci.edu/ml/datasets/Housing>

²Iris Data Set – <http://archive.ics.uci.edu/ml/datasets/Iris>



Example 12.3 Recognize handwritten digits.³



Each handwritten digit is preprocessed and transformed into a 28×28 greyscale image. Therefore, each input vector composes of 784 features.


³The MNIST Database of Handwritten Digits (<http://yann.lecun.com/exdb/mnist/>); Y. LeCun, L. Bottou, Y. Bengio and P. Haffner: Gradient-Based Learning Applied to Document Recognition, Proceedings of the IEEE, 86(11):2278-2324, November 1998

Example 12.4 Classify a set of artificially generated examples

x_1	x_2	t
2.124	-0.065	1
0.253	0.807	1
1.454	-0.578	1
0.569	0.573	1
0.458	-0.337	1
-0.809	0.297	1
0.864	-0.375	1
0.202	0.155	1
-1.875	0.705	1
-0.569	0.845	1

x_1	x_2	t
1.342	1.182	2
2.568	1.583	2
2.515	2.034	2
1.384	2.211	2
2.926	2.310	2
0.714	1.808	2
3.430	2.011	2
1.575	1.983	2
1.597	1.430	2
2.604	2.264	2

x_1	x_2	t
4.393	0.059	3
4.584	1.761	3
3.370	1.720	3
4.192	1.379	3
4.649	2.073	3
3.849	1.173	3
4.401	1.839	3
4.494	2.292	3
5.567	2.538	3
3.928	2.438	3



12.1.3 Probabilistic Prediction

Using the probability theory, we can find a prediction function from

$$y = \hat{f}(\mathbf{x}) = \underset{c=1}{\operatorname{argmax}}^C P(t = c | \mathbf{x}, \mathcal{D})$$

This is known as a **MAP estimate**.⁴


input
a training set
output
=0

12.2 K -Nearest Neighbors

A simple technique to predict the class of an unknown feature vector \mathbf{x} is to check the classes of the K nearest vectors in the training set. Then, predicting the class of \mathbf{x} using the majority class.

The distance between two vectors can be simply calculated by the Euclidean norm of the difference of the two vectors i.e.

$$\begin{aligned} \operatorname{dist}(\mathbf{x}_i, \mathbf{x}_j) &= \|\mathbf{x}_i - \mathbf{x}_j\| \\ &= \sqrt{(x_{i1} - x_{j1})^2 + (x_{i2} - x_{j2})^2 + \cdots + (x_{iD} - x_{jD})^2} \end{aligned}$$



$$\vec{x}_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{iD} \end{bmatrix} \quad \vec{x}_j = \begin{bmatrix} x_{j1} \\ x_{j2} \\ \vdots \\ x_{jD} \end{bmatrix}$$

Formally, we can compute the prediction probability from

$$P(t = c | \mathbf{x}, \mathcal{D}, K) = \frac{1}{K} \sum_{i \in N_K(\mathbf{x}, \mathcal{D})} \mathbb{I}(t_i = c)$$

where $N_K(\mathbf{x}, \mathcal{D})$ is the set of K nearest points to \mathbf{x} in \mathcal{D} , and $\mathbb{I}(e)$ is the indicator function defined as

$$\mathbb{I}(e) = \begin{cases} 1 & \text{if } e \text{ is true} \\ 0 & \text{if } e \text{ is false} \end{cases}$$

⁴MAP stands for *maximum a posteriori*

Exercise 12.1 From the following dataset, use the K -nearest neighbors technique to predict the classes of the following examples with different values of K .

x_1	x_2	t	x_1	x_2	t
0.5	0.5	0	2.0	1.5	1
0.2	1.0	0	3.0	1.0	1
1.0	0.8	0	4.0	3.2	1

1. $[0.0, 0.0]^T$ using $K = 3$

$$d\left(\begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \begin{bmatrix} 0.5 \\ 0.5 \end{bmatrix}\right) = \sqrt{(0.0-0.5)^2 + (0.0-0.5)^2} = \sqrt{0.5} \quad * \rightarrow 0$$

$$d\left(\begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \begin{bmatrix} 0.2 \\ 1.0 \end{bmatrix}\right) = \sqrt{(0.0-0.2)^2 + (0.0-1.0)^2} = \sqrt{1.04} \quad * \rightarrow 0$$

$$d\left(\begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \begin{bmatrix} 1.0 \\ 0.8 \end{bmatrix}\right) = \sqrt{(0.0-1.0)^2 + (0.0-0.8)^2} = \sqrt{1.64} \quad * \rightarrow 0$$

$$d\left(\begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \begin{bmatrix} 2.0 \\ 1.5 \end{bmatrix}\right) = \sqrt{(0.0-2.0)^2 + (0.0-1.5)^2} = \sqrt{6.25} \quad \text{Predict "0" for } \begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}$$

$$d\left(\begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \begin{bmatrix} 3.0 \\ 1.0 \end{bmatrix}\right) = \sqrt{(0.0-3.0)^2 + (0.0-1.0)^2} = \sqrt{10}$$

$$d\left(\begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \begin{bmatrix} 4.0 \\ 3.2 \end{bmatrix}\right) = \sqrt{(0.0-4.0)^2 + (0.0-3.2)^2} = \sqrt{26.24}$$

2. $[1.0, 1.0]^T$ using $K = 5$

$$d\left(\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}, \begin{bmatrix} 0.5 \\ 0.5 \end{bmatrix}\right) = \sqrt{0.5} \quad * \rightarrow 0$$

$$d\left(\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}, \begin{bmatrix} 0.2 \\ 1.0 \end{bmatrix}\right) = \sqrt{0.64} \quad * \rightarrow 0$$

$$d\left(\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}, \begin{bmatrix} 1.0 \\ 0.8 \end{bmatrix}\right) = \sqrt{0.04} \quad * \rightarrow 0$$

$$d\left(\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}, \begin{bmatrix} 2.0 \\ 1.5 \end{bmatrix}\right) = \sqrt{1.25} \quad * \rightarrow 1$$

$$d\left(\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}, \begin{bmatrix} 3.0 \\ 1.0 \end{bmatrix}\right) = \sqrt{4} \quad * \rightarrow 1$$

$$d\left(\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}, \begin{bmatrix} 4.0 \\ 3.2 \end{bmatrix}\right) = \sqrt{13.84}$$

Predict "0" for $\begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}$

input \rightarrow calculate distances to all training vectors \rightarrow find vectors with k smallest distances \downarrow voting

$$\begin{aligned} \text{input} &= \begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix} \\ d\left(\begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix}, \begin{bmatrix} 0.5 \\ 0.5 \end{bmatrix}\right) &= \sqrt{8.5} \\ d\left(\begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix}, \begin{bmatrix} 0.2 \\ 1.0 \end{bmatrix}\right) &= \sqrt{5.49} \\ d\left(\begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix}, \begin{bmatrix} 1.0 \\ 0.8 \end{bmatrix}\right) &= \sqrt{3.89} \\ d\left(\begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix}, \begin{bmatrix} 2.0 \\ 1.5 \end{bmatrix}\right) &= \sqrt{1} \\ d\left(\begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix}, \begin{bmatrix} 3.0 \\ 1.0 \end{bmatrix}\right) &= \sqrt{3.25} \\ d\left(\begin{bmatrix} 2.0 \\ 2.5 \end{bmatrix}, \begin{bmatrix} 4.0 \\ 3.2 \end{bmatrix}\right) &= \sqrt{9.25} \end{aligned}$$

12.3 Model Selection

Since we can run KNN using different value of K , which value of K should we use?

A basic way to evaluate a classifier is to use the *misclassification rate* on the training set. It can be defined as:

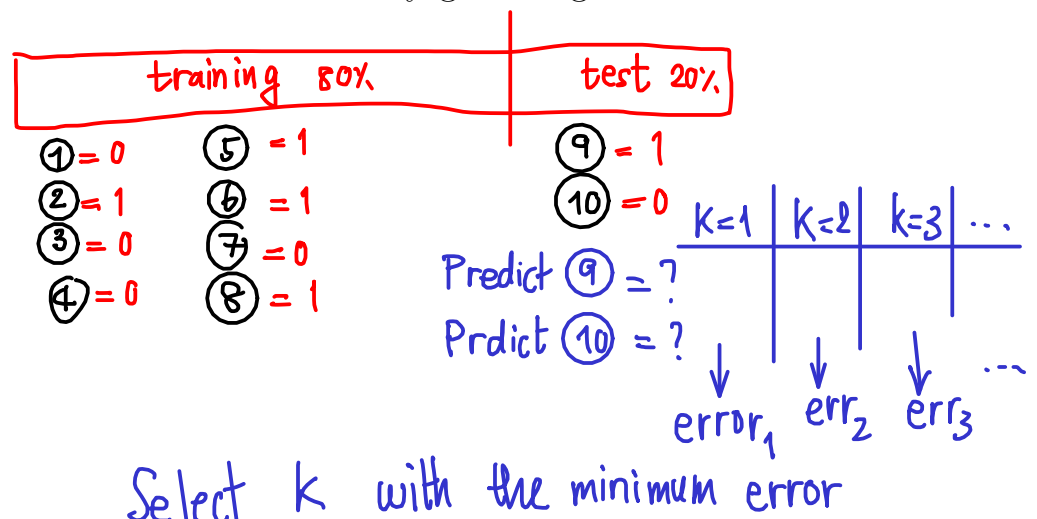
$$\text{err}(f, \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(f(\mathbf{x}_i) \neq t_i)$$

However, the error on the training set should not represent the *generalization error*.

12.3.1 Holdout Method


The collected dataset is split into two subsets, i.e. *training* and *validation* sets. Generally, about 80% of the data are for the training set, and 20% of the data are for the validation set.

However, the evaluation result depends heavily on the examples in the training set and the validation set. We may get too good or too bad result easily.




12.3.2 Cross Validation Method

The cross validation method splits the dataset into k folds. Then, it runs the holdout evaluation k runs. In each run, the k th fold is used as the validation set, while the combination of the rest folds are used as the training set in that run.




12.4 Artificial Neural Networks

Artificial Neural Network (ANN) is a supervised learning technique inspired by neurons in the nervous system.



Perceptron is a type of ANN working as a kind of binary classification based on a linear model. *two classes of output*



Each perceptron has several inputs expressing instance features denoting $(x_0, x_1, x_2, \dots, x_p)^T$. A weight is attached to each input as a vector $(w_0, w_1, w_2, \dots, w_p)^T$. These weights are the important part denoting the main characteristic of each perceptron. An output of perceptron denotes the prediction result.

A perceptron works as a **linear classifier** getting features as inputs and estimating the class. The prediction is done by

Given an input vector $(x_1, x_2, \dots, x_p)^T$, we have


$$\begin{aligned} \text{predicted output } \boxed{y} &= \text{sgn}\left(\sum_{i=0}^p w_i x_i\right) \\ &= \text{sgn}\left(w_0 + \underbrace{w_1 x_1 + \dots + w_p x_p}_{\text{weighted sum of the input vector}}\right) \end{aligned}$$

$$\text{where, } \text{sgn}(x) = \begin{cases} 1 & \text{if } x > 0, \\ 0 & \text{otherwise} \end{cases}$$

↓
the sign function

This classifier deals with a two-class problem (1 and 0). We have $w_0 + w_1 x_1 + \dots + w_p x_p = 0$ as a linear hyperplane.

$$\hat{y} = \text{sgn}(-2.44 + 0.5x_1 + 0.4x_2)$$




a perceptron = a line that classified the space into two areas (1, 0)
(with appropriate weights)

Exercise 12.2 Use the perceptron in the previous figure to predict the following examples.

1. $\begin{matrix} x_1 & x_2 \\ [0.0, 0.0]^T \end{matrix}$

$$\begin{aligned} \hat{y} &= \text{sgn}(-2.44 + (0.5)(0) + (0.4)(0)) \\ &= \text{sgn}(-2.44) = 0 \end{aligned}$$


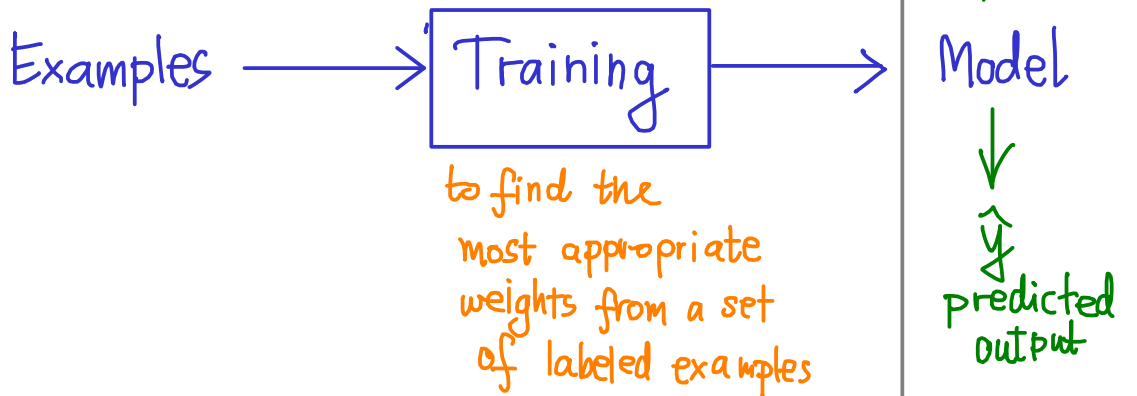


2. $[4.0, 4.0]^T$

$$\begin{aligned} \hat{y} &= \text{sgn}(-2.44 + (0.5)(4.0) + (0.4)(4.0)) \\ &= \text{sgn}(1.16) = 1 \end{aligned}$$

3. $[3.0, 2.0]^T$

$$\begin{aligned} \hat{y} &= \text{sgn}(-2.44 + (0.5)(3.0) + (0.4)(2.0)) \\ &= \text{sgn}(-0.14) = 0 \end{aligned}$$




Exercise 12.3 Answer the questions from the following dataset.

x_1	x_2	t
0	0	0
0	1	0
1	0	0
1	1	1


Training examples

1. Draw the structure of a perceptron for learning from this dataset.



We need to find an appropriate vector of weights from the given dataset.

2. Find the weights that makes the perceptron correctly classify the training examples.



From $(0, 1.5)$, we have $w_0 + w_1(0) + w_2(1.5) = 0$

$$w_0 + 1.5w_2 = 0 \quad \text{--- (1)}$$

From $(1.5, 0)$, we have $w_0 + w_1(1.5) + w_2(0) = 0$


$$w_0 + 1.5w_1 = 0 \quad \text{--- (2)}$$

$$(2) - (1) \quad 1.5w_1 - 1.5w_2 = 0 \quad \rightarrow \quad w_1 = w_2$$

when $w_2 = +1 \longrightarrow w_1 = +1$

substitute w_2 to (1) $\longrightarrow w_0 = -1.5$

$$-1.5 + x_1 + x_2 = 0$$




12.4.1 Perceptron Training Algorithm

Weights are the main component of each perceptron. If each weight is set to a suitable value, the perceptron will be able to classify unlabeled instances. We need a technique to assign or adjust weights until we obtain the most suitable values.

Obtaining the Weight Updating Rule

Gradient Descent technique can be used to find the most suitable weights. It is a simple technique to solve optimization problem on a continuous environment.

1. Define an objective function. In this case, it should be a function of \mathbf{w} i.e. $E(\mathbf{w})$. → error function
2. Randomly select initial values for \mathbf{w} .
3. Update weights using $\mathbf{w} - \eta \frac{dE}{d\mathbf{w}}$ where η is a positive small number.
4. Iteratively update until $\frac{dE}{d\mathbf{w}} \approx 0$



We can define an objective function using weights \mathbf{w} as a parameter. The function should return 0 if the weights are adjusted to the best values. In the simplest way, we can define an objective function as

$$E(\mathbf{w}) \stackrel{d}{=} \frac{1}{2} \sum_{d=1}^N (t_d - y_d)^2$$

$$= \frac{1}{2} \sum_{d=1}^N \left(t_d - \sum_{i=0}^p w_i x_{di} \right)$$

Error function
 target output
 predicted output
 a perceptron

We compute gradient of the objective function which is denoted as

$$\nabla E(\mathbf{w}) \stackrel{d}{=} \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_p} \right]$$

Then, we find the partial differentiate of E by w_i :

$$\begin{aligned}
 \frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \left(\frac{1}{2} \sum_{d=1}^N (t_d - y_d)^2 \right) \\
 &= \frac{1}{2} \sum_{d=1}^N \frac{\partial}{\partial w_i} \left((t_d - y_d)^2 \right) \\
 &= \frac{1}{2} \sum_{d=1}^N 2(t_d - y_d) \frac{\partial}{\partial w_i} (t_d - y_d) \\
 &= \sum_{d=1}^N (t_d - y_d) \frac{\partial}{\partial w_i} \left(t_d - \sum_{i=0}^p w_i x_{di} \right) \\
 &= \sum_{d=1}^N (t_d - y_d) (-x_{di}) = \nabla E(w_i)
 \end{aligned}$$

From the gradient of E , we have

$$w_i^* = w_i - \eta \nabla E(w_i)$$

a new weight w_i
the current w_i

We then update w_i by

$$w_i \leftarrow w_i + \eta(t - y)(x_i)$$

input
Weight updating rule
target output
predicted output

feed an input vector \longrightarrow make a prediction
 if there is a error \longrightarrow update weights
 ($t \neq y$)


→ find the most appropriate weights for a perceptron.

A training algorithm for perceptron is as below:

1. Start with random weights.

↗ $t \neq y$

2. Apply each training instance. If the perceptron misclassifies, then adjust all the weights. Each weight w_i for an input x_i is adjusted by


$$w_i \leftarrow w_i + \eta(t - y)x_i$$

where η is a positive constant called the *learning rate*.

3. Iteratively, perform step 2 through the training set until the perceptron classifies all training instances correctly.


Example 12.5 Find the most suitable weights from the following examples:

No	x_1	x_2	t	No	x_1	x_2	t
1	2.5	2.0	0	8	4.0	3.0	1
2	1.2	3.0	0	9	3.8	4.5	1
3	2.1	3.0	0	10	3.2	2.5	1
4	2.4	2.3	0	11	3.3	4.0	1
5	2.0	2.5	0	12	2.5	4.2	1
6	1.5	2.4	0	13	4.0	1.5	1
7	1.8	1.2	0	14	3.0	3.2	1

We start from random weights and the learning rate:

$$w_0 = -0.40; \quad w_1 = 0.18; \quad w_2 = 0.20; \quad \eta = 0.01$$

From this setting, we have a linear classifier as the following figure.




We start from applying an instance $([2.5, 2.0]^T, 0)$ to the perceptron, then

$$\begin{aligned}
 y &= \text{sgn}(0.18 \times 2.5 + 0.2 \times 2.0 - 0.4) \\
 &= \text{sgn}(0.45) \\
 &= \textcircled{1} \leftarrow \text{predicted output}
 \end{aligned}$$

Since $y \neq t$, then we adjust weight as follow:

$$\begin{aligned}
 w_0 &\leftarrow -0.40 + ((0.01)(0 - 1)(1.0)) = -0.410 \\
 w_1 &\leftarrow +0.18 + ((0.01)(0 - 1)(2.5)) = +0.155 \\
 w_2 &\leftarrow +0.20 + ((0.01)(0 - 1)(2.0)) = +0.180
 \end{aligned}$$


↓
current
 w_i
↓
 η
↓
 t
↓
 y
↓
 x_i
↓
new
 w_i



#2 ~ #14

After applying all the remaining 13 examples, we will get weights as:


$$w_0 = -0.440; \quad w_1 = 0.098; \quad w_2 = 0.097$$



#1 ~ #14 (5 times)

After updating the weights for five rounds, the weights become:


$$w_0 = -0.450; \quad w_1 = 0.091; \quad w_2 = 0.077$$



12.4.2 Threshold Function


A perceptron outputs either '0' or '1' based on the sum of the inputs and the weights, as well as the threshold function (e.g. $\text{sgn}(\cdot)$).

$$y = \text{sgn}(w_0 + w_1x_1 + w_2x_2)$$



The threshold function $\text{sgn}(\cdot)$ is a discontinuous function. It cannot be differentiated. Therefore, it cannot be used together with the gradient descent, or other optimization techniques.

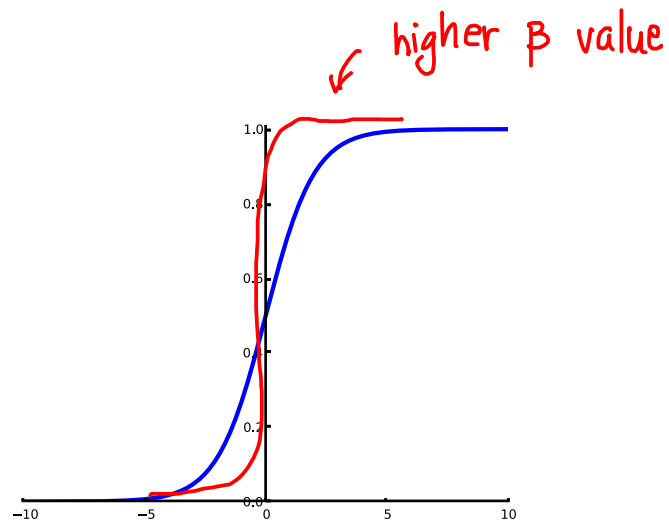
$$\text{sgn}(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$



The sigmoid function is an S-shaped function. It can be in place of the function $\text{sgn}(\cdot)$.

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-\beta x}}$$


$\beta = \text{a positive constant}$



12.4.3 Sigmoid Unit

Now, an output of a perceptron can be calculated from:

$$y = \text{sigmoid}(w_0 + w_1x_1 + w_2x_2)$$

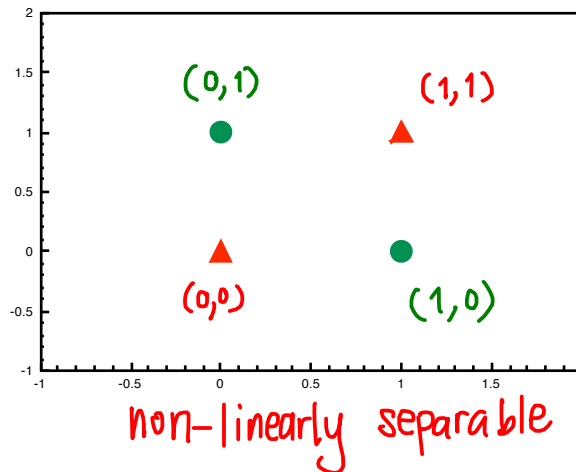


12.5 Linear Separability and Multilayer Perceptron

Since a perceptron represents a linear model for classification, it works very well on a training set that is **linearly separable**. However, a perceptron cannot deal with the training set that cannot be separated by just a line.

XOR Problem

x_1	x_2	t
0	0	0
0	1	1
1	0	1
1	1	0




a perceptron
 \equiv a line
 (a hyperplane)

How can the perceptron be improved to handle this situation?


To solve the non-linearly separable problem, we connect multiple sigmoid units to become a neural network.

The units are arranged into more than two layers. The technique is called multilayer perceptron. = a neural network




Each perceptron works as a linear discriminant. The inputs are classified by the linear discriminants in the hidden layer. Either 0 or 1 are yielded from each unit in the hidden layer. The outputs from the hidden layer are fed as inputs to the output layer. Here, the perceptron in the output is a linear discriminant that classifies the pre-classified inputs.

Example 12.6 We use two lines to correctly discriminate the examples:



These two lines can be represented as the following network:



$$z_1 = \text{sigmoid}(x_1 + x_2 - 0.5)$$


$$z_2 = \text{sigmoid}(x_1 + x_2 - 1.5)$$

From the network, we have two intermediate outputs (i.e. z_1 , and z_2) from the perceptrons in the hidden layer:


hidden layer

x_1	x_2	z_1	z_2	t
0	0	0	0	0
0	1	0	1	1
1	0	0	1	1
1	1	1	1	0

We can use two lines to correctly discriminate the examples:



Exercise 12.4 Use the multilayer perceptron with the following weights to predict the output of $[1.0, 1.0]^T$. Set $\beta = 1$.



$$\begin{aligned} w_{01}^H &= +0.50, & w_{11}^H &= -0.50, & w_{21}^H &= +0.10, \\ w_{02}^H &= +1.00, & w_{12}^H &= +0.20, & w_{22}^H &= -0.20, \\ w_{03}^H &= -1.00, & w_{13}^H &= +0.10, & w_{23}^H &= +0.50, \\ w_{01}^Y &= -1.00, & w_{11}^Y &= +1.50, & w_{21}^Y &= -1.00, & w_{31}^Y &= 1.00 \end{aligned}$$

$$\begin{aligned} z_1 &= \text{sigmoid}(w_{01}^H + w_{11}^H x_1 + w_{21}^H x_2) = \text{sigmoid}(0.5 + (-0.5)(1.0) + (0.1)(1.0)) \\ &= \text{sigmoid}(0.1) = \frac{1}{1 + e^{-(1)(0.1)}} = 0.52 \end{aligned}$$

$$\begin{aligned} z_2 &= \text{sigmoid}(w_{02}^H + w_{12}^H x_1 + w_{22}^H x_2) = \text{sigmoid}(1 + (0.2)(1) + (-0.2)(1)) \\ &= \text{sigmoid}(1) = 0.73 \end{aligned}$$


$$\begin{aligned} z_3 &= \text{sigmoid}(w_{03}^H + w_{13}^H x_1 + w_{23}^H x_2) = \text{sigmoid}(-1 + (0.1)(1) + (0.5)(1)) \\ &= \text{sigmoid}(-0.4) = 0.40 \end{aligned}$$

$$\begin{aligned} y &= \text{sigmoid}(w_{01}^Y + w_{11}^Y z_1 + w_{21}^Y z_2 + w_{31}^Y z_3) \\ &= \text{sigmoid}(-1 + (1.5)(0.52) + (-1)(0.73) + (1)(0.40)) \\ &= \text{sigmoid}(-0.58) = 0.36 \end{aligned}$$

using the threshold of 0.5,

the predicted output is 0

$$(0.36 < 0.5)$$



12.5.1 Backpropagation Algorithm

The gradient descent technique is used to train a multilayer perceptron in the same manner as training a perceptron.

We first define the error function:

$$E(\mathbf{w}^{\mathcal{H}}, \mathbf{w}^{\mathcal{Y}}) = \frac{1}{2} \sum_{d=1}^N (t_d - y_d)^2$$

We can use the chain rule to calculate the gradient

$$\frac{\partial E}{\partial w_{jl}^{\mathcal{H}}} = \frac{\partial E}{\partial y} \frac{\partial y}{\partial z_l} \frac{\partial z_l}{\partial w_{jl}^{\mathcal{H}}}$$

Here, the error propagates from the output back to the inputs. Thus, the algorithm is named *backpropagation*.

For each example in D :

Forward phase: Given an input vector, compute the predicted output

1. compute the output of each perceptron in the hidden layer:

$$z_l = \text{sigmoid}\left(\sum_i w_{il}^{\mathcal{H}} x_i\right)$$

2. compute the output of each perceptron in the output layer:

$$y_k = \text{sigmoid}\left(\sum_l w_{lk}^{\mathcal{Y}} z_l\right)$$

Backward phase: Update $\vec{w}^{\mathcal{H}}$ and $\vec{w}^{\mathcal{Y}}$ according to the error $(t - y)$

1. compute the error at the output using:

$$\delta_k^{\mathcal{Y}} = (t_k - y_k) y_k (1 - y_k)$$

2. update the output layer weights using:

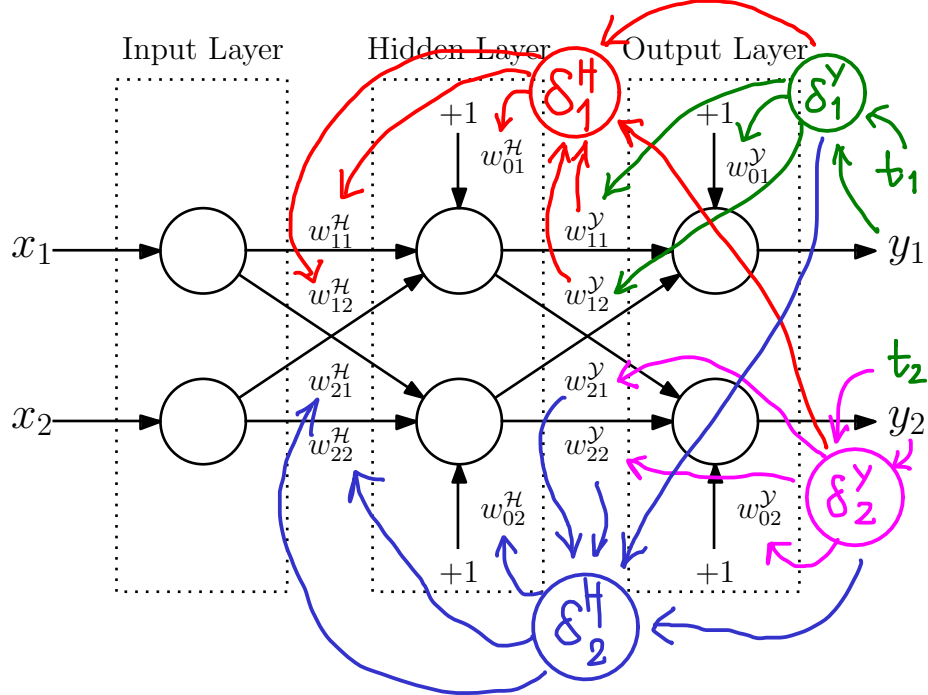
$$w_{lk}^{\mathcal{Y}} \leftarrow w_{lk}^{\mathcal{Y}} + \eta \delta_k^{\mathcal{Y}} z_l$$

3. compute the error in the hidden layer using:


$$\delta_l^{\mathcal{H}} = z_l (1 - z_l) \sum_k w_{lk}^{\mathcal{Y}} \delta_k^{\mathcal{Y}}$$

4. update the hidden layer weights using:

$$w_{il}^{\mathcal{H}} \leftarrow w_{il}^{\mathcal{H}} + \eta \delta_l^{\mathcal{H}} x_i$$



Exercise 12.5 Given the following multilayer perceptron with all initial weights of 0's Find the all updated values of the weights of all layers when we feed a labeled example $([1.0, 0.5]^T, 1)$. Set $\beta = 1$ and $\eta = 0.01$.



Backward phase

$$\delta_1^y = (t_1 - y_1)(y_1)(1 - y_1) = 0.125$$

$$w_{11}^y = 0 + \eta \delta_1^y z_1 = 0 + (0.01)(0.125)(0.5) = 0.000625$$

$$w_{21}^y = 0 + \eta \delta_1^y z_2 = 0.000625$$

$$w_{31}^y = 0 + \eta \delta_1^y z_3 = 0.000625$$

$$w_{01}^y = 0 + \eta \delta_1^y (1) = 0.00125$$

$$\begin{aligned} \delta_1^H &= z_1(1 - z_1) \sum_k w_{1k}^y \delta_k^y \\ &= (0.5)(1 - 0.5)(0.000625)(0.125) = 1.95 \times 10^{-5} \end{aligned}$$

(k) \rightarrow #output units

$$w_{01}^H = 0 + \eta \delta_1^H (1) = 0 + (0.01)(1.95 \times 10^{-5})(1) = 1.95 \times 10^{-7}$$

$$w_{11}^H = 0 + \eta \delta_1^H x_1 = 0 + (0.01)(1.95 \times 10^{-5})(1) = 1.95 \times 10^{-7}$$

$$w_{21}^H = 0 + \eta \delta_1^H x_2 = 0 + (0.01)(1.95 \times 10^{-5})(0.5) = 9.75 \times 10^{-8}$$

$$\delta_2^H \rightarrow w_{02}^H, w_{12}^H, w_{22}^H$$

$$\delta_3^H \rightarrow w_{03}^H, w_{13}^H, w_{23}^H$$

References

- Satish Kumar, “Neural networks: a classroom approach”, McGraw-Hill, 2005.
- Stephen Marsland, “Machine Learning: an algorithm perspective”, CRC Press, 2009.
- Ethem Alpaydin, “Introduction to Machine Learning”, The MIT Press, 2004.
- Kevin P. Murphy, “Machine Learning: A Probabilistic Perspective”, The MIS Press, 2012.